

# The Java Concurrency API

**Cristian Mateos<sup>1,2</sup>**

<sup>1</sup>ISISTAN-UNICEN-CONICET

<sup>2</sup>Consejo Nacional de Investigaciones Científicas y Técnicas (CONICET)

Computación Paralela y Distribuida

# Agenda

- The Java Concurrency API: Why
- Java Concurrency API: Basics
  - Executors: thread pools and scheduling
  - Futures
  - Synchronizers
  - Atomic variables and concurrent collections
- The ForkJoin framework

## Why concurrency utilities

- Recall that Java has built-in concurrency primitives - *wait()*, *notify()*, and *synchronized*

... but they all have shortcomings

- Hard to use correctly; easy to use incorrectly
- Too low level for many applications
- Can lead to poor performance if used incorrectly
- Leave out lots of useful concurrency constructs

# Why concurrency utilities

- Recall that Java has built-in concurrency primitives - *wait()*, *notify()*, and *synchronized*

## ... but they all have shortcomings

- Hard to use correctly; easy to use incorrectly
- Too low level for many applications
- Can lead to poor performance if used incorrectly
- Leave out lots of useful concurrency constructs

## Concurrency utilities: Goals

On one hand, provide efficient, correct and reusable concurrency building blocks. On the other hand, enhance scalability, performance, readability, maintainability, and thread-safety of concurrent Java applications.



# Building blocks

- Executor, ThreadPool, and Future
- Synchronizers: Semaphores and Barriers
- Atomic Variables
  - Low-level compare-and-set operations
- Concurrent collections
  - *BlockingQueue*, *SynchronousQueue*, *ConcurrentHashMap*
- Recommended reading: “Java Concurrency in Practice” by Brian Goetz et al.

# Executor

- Standardizes asynchronous invocation
- Separates job submission from execution policy
  - *anExecutor.execute(aRunnable)*
  - not *new Thread(aRunnable).start()*
- Two code styles supported:
  - Actions: *Runnable*s
  - Functions: *Callable*s
  - Also has lifecycle management: e.g., cancellation, shutdown
- Executor usually created via Executors factory class
  - Implicitly associated to a *ThreadPoolExecutor*
  - Via *ExecutorService*, customizes shutdown methods, before/after hooks, saturation policies, queuing

## Executor and ExecutorService

- ExecutorService adds lifecycle management to Executor

```
public interface Executor {  
    void execute(Runnable command);  
}  
public interface ExecutorService extends Executor {  
    void shutdown();  
    List<Runnable> shutdownNow();  
    boolean isShutdown();  
    boolean isTerminated();  
    boolean awaitTermination(long timeout, TimeUnit unit);  
    // other convenience methods for submitting tasks  
}
```



## Executors factory methods

```
public class Executors {  
    static ExecutorService newSingleThreadedExecutor();  
    static ExecutorService newFixedThreadPool(int n);  
    // Beware demand peaks and long-running threads!  
    static ExecutorService newCachedThreadPool(int n);  
    static ScheduledExecutorService  
        newScheduledThreadPool(int n);  
    // additional versions & utility methods  
}
```

## Example without Executors: Web Server

- 1 Thread per message received (no limit on thread creation)

```
class WebServer {  
    public static void main(String[] args) {  
        ServerSocket socket = new ServerSocket(80);  
        while (true) {  
            final Socket connection = socket.accept();  
            Runnable r = new Runnable () {  
                public void run(){handleRequest(connection);}  
            };  
            new Thread(r).start();  
        }  
    }  
    public static void handleRequest(Socket socket){...}  
}
```

## Example with Executors: Web Server

- Thread pool-based Web Server... better resource management!

```
class WebServer {  
    Executor pool = Executors.newFixedThreadPool(7);  
    public static void main(String[] args) {  
        ServerSocket socket = new ServerSocket(80);  
        while (true) {  
            final Socket connection = socket.accept();  
            Runnable r = new Runnable() {  
                public void run() {handleRequest(connection);}  
            };  
            pool.execute(r);  
        }  
    }  
    public static void handleRequest(Socket socket){...}  
}
```

## Callable and Future

- *Callable* is functionally analog of *Runnable*

```
interface Callable<V> {  
    V call() throws Exception;  
}
```

- *Future* represents result of asynchronous computation

```
interface Future<V> {  
    V get() throws InterruptedException , ExecutionException;  
    V get(long timeout, TimeUnit unit);  
    boolean cancel(boolean mayInterruptIfRunning);  
    boolean isCancelled();  
    boolean isDone();  
}
```

# Using Futures

- Client initiates asynchronous computation via one-way message
- Client receives a “handle” to the result: a *Future*
- Client does other work while waiting for result
- When ready, Client requests result from *Future*, blocking if necessary until result is available
- Client uses result
- **Example:** Searching a big file for a given string

# File searcher

*// I am consciously ignoring exception handling...*

```
public class FutureFileSearcher {  
    ExecutorService executor = Executors.newFixedThreadPool(1);  
    String searchFile(String keyword, String filePath) {  
        Future<String> future = executor.submit(  
            new Callable<String>() {  
                public String call() {  
                    return doFileSearch(keyword, filePath);  
                }  
            });  
        // Do some other useful stuff here  
        // Check every 0.5 secs  
        // (just for the sake of exemplification!)  
        while (!future.isDone())  
            Thread.sleep(500);  
        return future.get();  
    }  
}
```

# FutureTask

- A cancellable asynchronous computation
- A base implementation of *Future*
- Can wrap a *Callable* or *Runnable*
  - Allows *FutureTask* to be submitted to an *Executor*

```
...
FutureTask<String> futureTask = new FutureTask<String>(
    new Callable<String>(){
        public String call() {
            return doFileSearch(keyword, filePath);
        }
    });
executor.execute(futureTask);
...
```

## FutureTask and conventional threads

```
FutureTask<Integer> task = new FutureTask<Integer>(() -> {  
    System.out.println("Something_complicated_is_computed");  
    Thread.sleep(1000);  
    return 42;  
});  
Thread t1 = new Thread(() -> {  
    try {  
        int r = task.get();  
        System.out.println("Result_is_" + r);  
    } catch (InterruptedException | ExecutionException e) {}  
});  
System.out.println("t1_wil_wait_until_computation_is_ready");  
t1.start();  
task.run();
```



## Future and FutureTask example: Cache

```
public class Cache<K, V> {  
    Map<K, Future<V>> map = new ConcurrentHashMap<K, Future<V>>();  
    Executor executor = Executors.newFixedThreadPool(8);  
    public V get (final K key) {  
        Future<V> f = map.get(key); // null if key not found  
        if (f == null) {  
            Callable<V> c = new Callable<V>() {  
                public V call() {// compute value associated with key}  
            };  
            f = new FutureTask<V>(c);  
            // if key not found put(key,f) & return null  
            Future old = map.putIfAbsent(key, f);  
            if (old == null) { // otherwise return get(key)  
                executor.execute(f);  
            } else { f = old; }  
        }  
        return f.get();  
    }  
}
```

## ScheduledExecutorService

- For deferred and recurring tasks, can schedule
  - *Callable* or *Runnable* to run once with a fixed delay after submission
  - Schedule a *Runnable* to run periodically at a fixed rate
  - Schedule a *Runnable* to run periodically with a fixed delay between executions
- Submission returns a *ScheduledFutureTask* handle which can be used to cancel the task
- Like *Timer*, but supports more threads, pooling and finer thread control, so it is more versatile

# Synchronizers

- *Semaphore*: Dijkstra counting semaphore, managing some number of permits
- *CountDownLatch*: allows one or more threads to wait for a set of threads to complete an action
- *CyclicBarrier*: allows a set of threads to wait until they all reach a specified barrier point
- *Exchanger*: allows two threads to rendezvous and exchange data, such as exchanging an empty buffer for a full one

# Semaphores

- Maintain a logical set of permits
- *acquire()* blocks until a permit is free, then takes it
- *release()* adds a permit, releasing a blocking acquirer
- Often used to restrict the number of threads that can access some resource
  - But can be used to implement many synchronization kernels, patterns and primitives

# CountDownLatch

- Latching variables are conditions that once set never change
- Often used to start several threads, while having them wait for a signal before continuing

Main :

```
CountDownLatch startSignal = new CountDownLatch(1);
CountDownLatch stopSignal = new CountDownLatch(COUNT);
ExecutorService executor = Executors.newFixedThreadPool(COUNT);
for (int i = 0; i < COUNT; i++){
    executor.execute(new Worker(startSignal, stopSignal, i));
}
executor.shutdown();
// Do other stuff here...
startSignal.countDown();
try {
    stopSignal.await();
} catch (InterruptedException ex) {
    ex.printStackTrace();
}
```

## CountDownLatch (cont.)

Worker:

```
public void run() {  
    try {  
        // wait until the latch has counted down to zero  
        startSignal.await();  
    } catch (InterruptedException ex) {  
        ex.printStackTrace();  
    }  
    System.out.println("Running:_" + name);  
    stopSignal.countDown();  
}
```

## Cyclic barrier

- Allows threads to wait at a common barrier point
- Useful when a fixed-sized party of threads must occasionally wait for each other
- Can be re-used after threads released
- Can execute a *Runnable* once per barrier point
  - After the last thread arrives, but before any are released
  - Useful for updating shared-state before threads continue

## Cyclic barrier: Example

```
final CyclicBarrier barrier = new CyclicBarrier(3);
ExecutorService executor = Executors.newFixedThreadPool(3);
for (int i = 0; i < 3; i++) {
    executor.execute(new Runnable(){
        public void run(){
            try {
                log("At_run()"); barrier.await();
                log("Do_work"); barrier.await();
                log("Wait_for_end"); barrier.await();
                log("Done");
            }
            catch (Exception e) { e.printStackTrace(); }
        }
    });
}
executor.shutdown();
```



# Exchanger

- Synchronization point where two threads exchange objects
- A bidirectional *SynchronousQueue*
- Each thread presents some object on entry to the *exchange()* method, and receives the object presented by the other thread on return
- Example: Two threads filling and emptying two different buffers

## Exchanger (cont.)

- `exchange(V x)`: It waits for another thread to arrive at exchange point and exchange object with that thread.
- `exchange(V x, long timeout, TimeUnit unit)`: It waits for another thread for specific time interval provided in the method and exchange object with that thread.

